A Systematic Literature Review of Large Language Models and Small Language Models in Hybrid, Multi-Agent, and Monolithic Architectures

Doğukan Topçu[1], Abdulbaki Çeltik[1], Berkhan Kargılı[1], Meriç Kelemeden[1], Onur Demirörs[2], and Hüseyin Ünlü[2]

[1] Izmir Institute of Technology, Izmir, Turkey
dogukantopcu@std.iyte.edu.tr
abdulbakiceltik@std.iyte.edu.tr
berkhankargili@std.iyte.edu.tr
merickelemeden@std.iyte.edu.tr
[2] Izmir Institute of Technology, Izmir, Turkey
Advisors
onurdemirors@iyte.edu.tr
huseyinunlu@iyte.edu.tr

Abstract. Large language models (LLMs) now define the state of the art in many reasoning, generation, and decision-support tasks, but their deployment cost, energy demand, latency, and hardware requirements remain substantial. Small language models (SLMs) provide an attractive alternative for edge deployment, privacy-sensitive settings, and task-specific optimisation, yet their capabilities are uneven across domains. This paper presents a systematic literature review of 54 studies on LLMs, SLMs, and SLM–LLM collaboration, with emphasis on hybrid, multi-agent, retrieval-centered, and monolithic architectures. The review is guided by four research questions addressing comparative performance, computational efficiency, orchestration mechanisms, and domain suitability. Using a PRISMA 2020-guided review design and a Web of Science-centered search process, we synthesize evidence on model choices, architectural patterns, datasets, metrics, baselines, latency, cost, energy, and reported limitations. The reanalysis of the corpus shows a strong shift toward hybrid and multi-agent designs together with a clear reporting asymmetry in which accuracy is common but cost and energy are rarely quantified. The findings indicate that SLMs are most competitive when they are paired with routing, retrieval, planning, or fusion mechanisms and when the problem space is narrow, structured, or domain-specific.

Keywords: systematic literature review · large language models · small language models · hybrid architectures · multi-agent systems · edge-cloud inference

1 Introduction

Large language models (LLMs) excel at reasoning, coding, and multimodal tasks [?,?,?]. However, due to the high hardware requirements (RAM, GPU, etc.) they demand, these models are usually run only on cloud-based infrastructures or very advanced local systems [?,?]. This creates a significant bottleneck for system developers seeking to reduce cloud dependency, ensure data privacy, and avoid high operational costs [?,?].

As a result of these limitations, small language models (SLMs) are emerging as a strong alternative. Their ability to run on local machines provides a significant advantage in terms of data security and deployment autonomy, while also offering substantially lower operational costs than cloud-scale LLM setups [?,?]. Moreover, recent studies show that SLMs can be adapted to specific domains via fine-tuning or distillation, enabling competitive performance in specialized tasks [?,?,?]. In practical systems, SLMs are also attractive because they integrate well with explicit control modules (e.g., rule-based filters, routing/verification thresholds, or planners), improving controllability and reliability beyond prompt-only solutions [?,?,?]. However, SLMs still struggle to match the broad expertise and robust reasoning of massive frontier models, especially on tasks requiring deep contextual understanding and extensive world knowledge [?,?]. Therefore, framing the problem as a simple "SLM or LLM?" competition is no longer sufficient; instead, current research investigates when a single SLM is adequate, when an LLM is necessary, and when hybrid or agent-based architectures are preferable [?,?,?].

The existing literature provides concrete, domain-specific examples showing that the choice between SLMs, LLMs, and hybrid designs is primarily a systems decision. In telecommunications, Symbiotic Agents reports up to a fivefold reduction in decision error and near-real-time operation (82 ms loops) with up to 99.9% lower GPU overhead using SLM variants [?]. In medicine, self-reflective pipelines pair a compact generator with an LLM-based judge to iteratively refine outputs and improve diagnostic reliability [?].

These developments provide practical alternatives for scenarios where relying on a single monolithic LLM is either insufficient for task demands or impractical under strict resource constraints. Accordingly, recent work increasingly targets measurable improvements beyond raw accuracy, including reduced computational load, lower operational cost, and tighter latency budgets [?,?,?,?]. Across domains, researchers have reported which hybrid design patterns are effective for specific task classes—for example, selective edge–cloud verification, expert routing, and retrieval- or planner-centered pipelines—and how these choices translate to real-world deployments [?,?,?,?].

Motivated by these trends, this review analyzes when a single SLM is sufficient, when an LLM remains necessary, and when hybrid or multi-agent architectures provide the best trade-off between capability and efficiency [?,?,?,?]. We further examine domain conditions under which fine-tuned SLMs can realistically replace or complement larger models, and where they still fall short, synthesizing recurring architectural motifs across telecom, medical, and cybersecurity settings [?,?,?].

The main contributions of this review are as follows: (1) It provides a synthesized overview of 56 studies covering cloud-based monolithic LLMs, hybrid SLM–LLM architectures, the performance of fine-tuned SLMs versus large models, and multi-agent workflows. (2) It reframes the literature around four distinct practical research questions: Performance, Efficiency, Architecture, and Domain Suitability. (3) It creates an analytical review matrix that reveals the commonalities and differences among the studies, mapping which specific research question each study addresses.

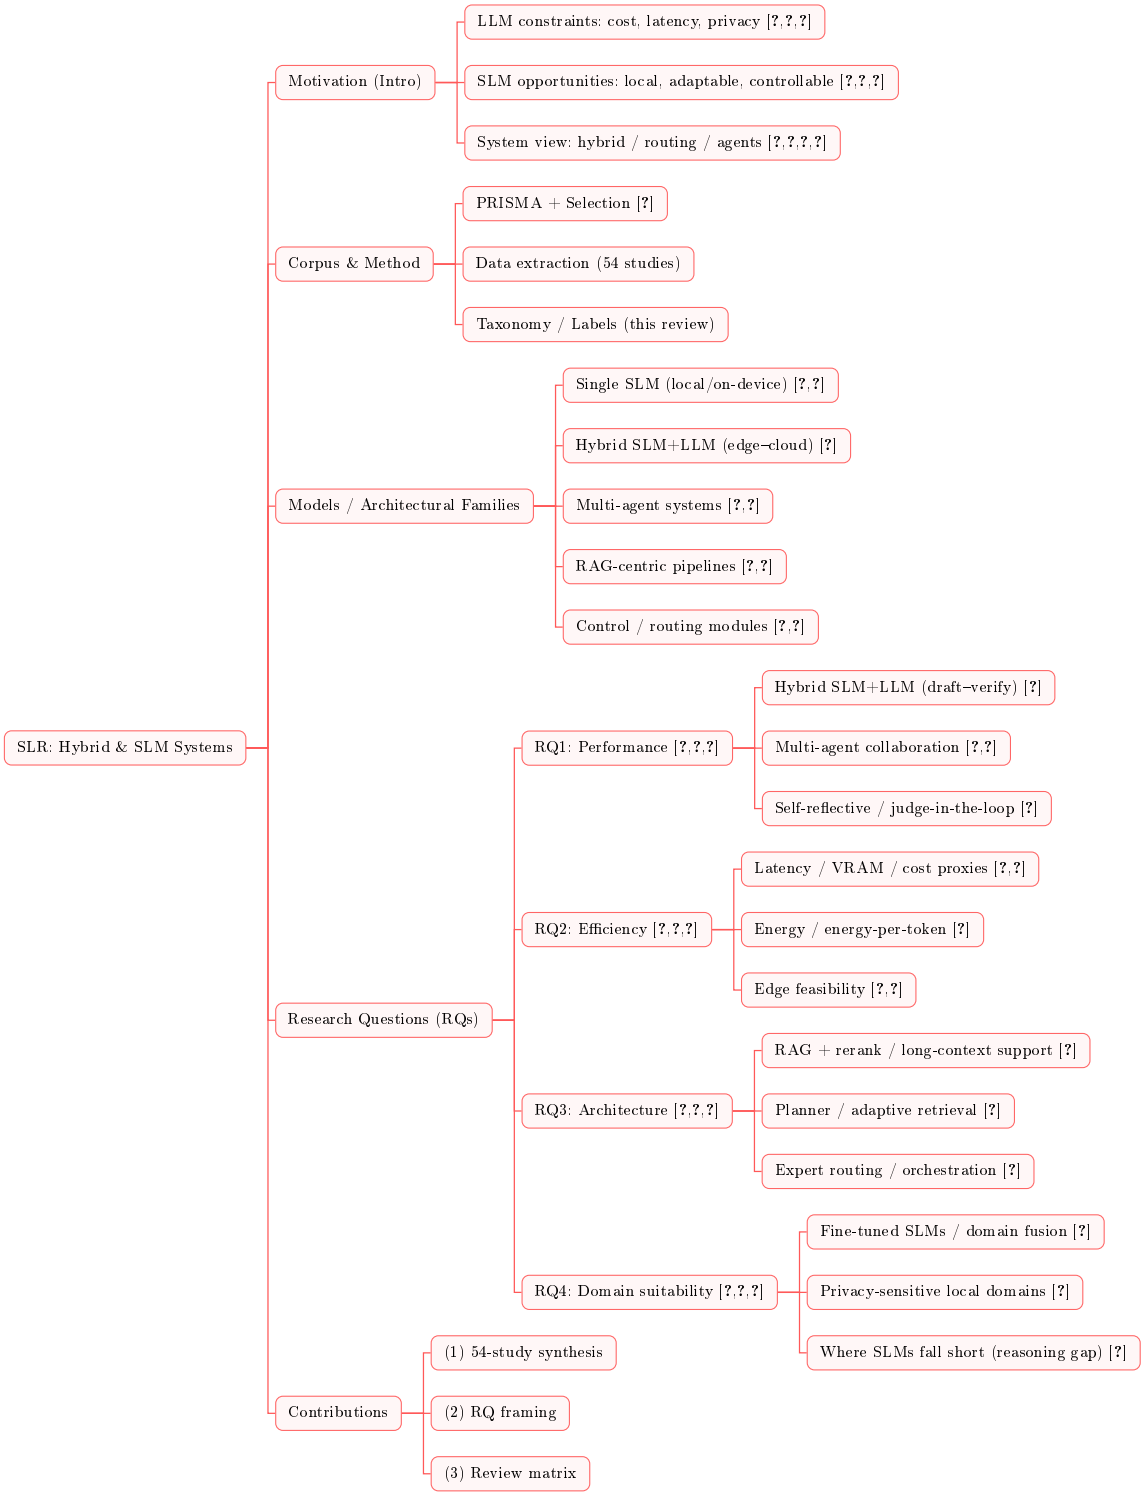

Fig. 1. Overview of the review structure and taxonomy used in this SLR (synthesized from the reviewed corpus; reporting guided by PRISMA [?]).

2 Related Works

2.1 Existing Systematic Literature Reviews and Surveys

The project website identifies five especially relevant prior reviews. The closest methodological competitor is the PRISMA-based SLR on small language models available on the project website. It is the nearest precursor because it combines a transparent review process with a broad overview of SLM research. However, it remains broader than the present study: its main focus is the SLM landscape itself, whereas our review centers on the relationship between SLMs, LLMs, and hybrid architectures.

The remaining relevant works are mostly narrative or semi-systematic surveys. *A Comprehensive Survey of Small Language Models in the Era of LLMs* provides wide coverage of definitions, training methods, applications, collaboration, and trustworthiness, but it does not follow a reproducible selection protocol. *A Survey on Collaborative Mechanisms Between Large and Small Language Models* offers a technically valuable taxonomy of routing, auxiliary, fusion, and distillation strategies, especially for edge deployment, but its scope is mechanism-centered rather than corpus-centered. *A Survey on Collaborating Small and Large Language Models for Performance, Cost-effectiveness, Cloud-edge Privacy, and Trustworthiness* overlaps substantially with the current review's themes, yet its methodology is still narrative rather than PRISMA-guided. Finally, *Collaborative Inference and Learning between Edge SLMs and Cloud LLMs* provides systematic organization of post-2023 edge-cloud work, but it is intentionally narrow in deployment context.

What emerges from these prior reviews is not a lack of attention to SLMs or collaboration, but a difference in analytical emphasis. Previous surveys often answer the question "What kinds of SLMs or collaboration mechanisms exist?" The present review instead asks "What patterns actually recur across recent primary studies, how do they affect performance and efficiency, and which domains benefit most from hybridization?" This difference is important because the research problem has shifted from model cataloguing to system design.

2.2 Main Gaps in the Existing Review Literature

The review website's comparative analysis surfaces four recurring gaps, all of which are visible in the primary studies as well.

First, application domains are underemphasized relative to architecture taxonomies. Prior surveys are strong on categories such as routing, distillation, and cloud-edge cooperation, but weaker on a systematic comparison of concrete domains such as telecom question answering, radiology reporting, optimisation modelling, dialogue state tracking, or industrial video captioning. This matters because whether an SLM is "good enough" depends heavily on task structure and domain constraints.

Second, most related works are narrative surveys rather than transparent systematic reviews. Narrative surveys are useful for breadth and conceptual framing, but they make it harder to reproduce study selection, audit exclusions, and compare cross-study evidence using a consistent extraction schema.

Third, prior reviews often stop before the newest wave of 2025 and early-2026 papers that focus on on-device agents, uncertainty-aware inference, prompt optimisation, multimodal collaboration, and energy-aware routing. Because this area evolves rapidly, even a one-year gap meaningfully changes the architectural picture.

Fourth, quantitative synthesis remains thin. Prior surveys usually discuss performance, efficiency, and privacy conceptually, but they rarely align cross-study indicators such as latency in milliseconds, API cost, FLOP reduction, or energy-per-token. This limits their practical usefulness for system builders who need to make deployment decisions rather than only understand the research landscape.

2.3 Positioning of This Study

This review is positioned as a methodologically explicit and engineering-oriented synthesis of SLM–LLM systems. It differs from prior work in three ways. First, it treats domain suitability as a first-class analytical axis rather than a side effect of architectural categories. Second, it explicitly maps research questions to extraction fields so that claims about performance, efficiency, and orchestration remain traceable to the underlying evidence. Third, it combines broad corpus-level analysis with representative deep dives into primary papers, which is necessary because the dataset's "key findings" fields are valuable for synthesis but too compressed to fully explain why a design works.

The website's own framing is also instructive: it argues that the review should be seen as a PRISMA-guided study of domain-specific applications of SLM–LLM hybrid architectures that emphasizes quantitative performance metrics. Our reanalysis supports this framing. In the current corpus, the most consequential papers are not simply "about small models" or "about collaboration" in the abstract; they are about domain-grounded deployment strategies under resource constraints.

3 Research Methodology

3.1 Review Protocol

The study follows a PRISMA 2020-guided review design [**?**]. The workflow used on the review website separates the process into identification, screening, and inclusion, with explicit reference to a Web of Science-centered search strategy. In the identification phase, query construction combines the key concepts of "small language models", "large language models", "hybrid architectures", "multi-agent systems", "orchestration", and "collaboration" using Boolean operators. In the screening phase, records are deduplicated, reviewed at title and abstract level, and then assessed at full text. In the inclusion phase, the selected studies are coded into a structured extraction sheet and then used to answer the research questions.

The website also makes an important methodological choice explicit: the review is not organized as a generic survey of all SLM work, but as a filtered corpus intended to support a comparison between monolithic, hybrid, and agentic designs. This matters because it shapes both what is included and how the findings are interpreted. Studies are not only selected because they use small models; they are selected because they help answer questions about comparative performance, orchestration, and efficiency.

3.2 Inclusion and Exclusion Criteria

The inclusion criteria extracted from the website can be summarized as follows. A study is included if it examines SLMs alone or SLM + LLM architectures, fits the project's small-model definition, focuses on at least one relevant technical question such as performance or orchestration, falls inside the contemporary review window, and provides empirical evaluation. These criteria are narrow

Table 1. Corpus overview derived from the 54-paper website snapshot.

Aspect	Observation
Corpus size	54 reviewed papers.
Publication trend	2025 dominates the corpus (38 papers), followed by 2024 (11 papers). A very small remainder falls in 2023, 2026, or metadata-anomalous records.
Architecture split	Hybrid: 18, Single-model: 14, Multi-agent: 12, Ensemble: 6, RAG-centered: 4.
Most frequent SLM families	Llama, Qwen, BERT-family models, Phi, Mistral, T5, and Gemma appear most often in the extracted studies.
Most frequent LLM families	GPT-4/GPT-4o, Llama, Qwen, GPT-3.5, Claude, and Gemini dominate the LLM side of the corpus.
Reporting coverage	Accuracy/task quality: 51 papers; latency: 17; cost: 5; energy: 4.
Recurring domains	Telecom QA, healthcare and radiology, mathematical reasoning, code and software tasks, dialogue systems, education, scientific prediction, and multimodal industrial monitoring.

enough to keep the corpus technically coherent while broad enough to admit retrieval-centered, multi-agent, and fusion-based systems.

The exclusion criteria remove five types of work: purely large-model comparisons with no SLM component, closed-box studies with insufficient transparency, non-technical ethical or regulatory discussions with no technical validation, studies outside the review time window, and purely theoretical work without experiments. Together, these choices bias the corpus toward deployable systems and measured outcomes, which is appropriate for the present research questions.

3.3 Corpus Profile

The current website snapshot contains 54 reviewed papers. The corpus is strongly concentrated in 2025, with 38 studies from that year, 11 from 2024, 2 from 2026, and a very small remainder in 2023 or metadata-anomalous records.

Architecturally, the corpus is not dominated by one single paradigm. Using the website's own filtering logic, 18 papers can be grouped as hybrid, 14 as single-model systems, 12 as multi-agent systems, 6 as ensembles, and 4 as RAG-centered systems. This matters because it shows that the field is not converging on a single "best" way to combine small and large models. Instead, the literature currently explores multiple control granularities, ranging from token-level speculative verification to query-level routing, sub-task decomposition, and representation fusion.

The corpus also reveals a strong reporting asymmetry. Accuracy or task-quality measures are reported in 51 of the 54 studies. In contrast, latency is reported in only 17, cost in only 5, and energy in only 4. This imbalance is central to interpreting the literature. Many papers claim efficiency benefits, but only a small subset actually report the variables needed to compare those claims rigorously across studies. Table **??** summarizes the corpus profile used in the review.

3.4 Data Extraction and Synthesis Strategy

For each study, the website records the SLM and LLM models used, architecture type, orchestration method, agent count, task type, dataset, evaluation metrics, baselines, accuracy, latency, cost, energy, key findings, and limitations. These fields are sufficient for a structured narrative synthesis because they capture both the functional design of the system and the evidence used to justify it. Table **??** maps those extraction fields to the four research questions.

The present manuscript uses that matrix in two steps. First, it performs corpus-level synthesis based on counts, categories, and recurring patterns across all 54 studies. Second, it supplements the

Table 2. RQ × data extraction matrix used in the review. A check mark indicates that the field directly supports the corresponding research question.

Extraction Field	RQ1	RQ2	RQ3	RQ4
Models Used (SLM)	✓			✓
Models Used (LLM)	✓			✓
Architecture	✓			✓
Orchestration Method			✓	✓
Agent Count		✓		✓
Task Type	✓		✓	
Dataset	✓		✓	
Metrics	✓		✓	✓
LLM Baseline	✓		✓	✓
SLM Baseline	✓		✓	✓
Accuracy	✓		✓	
Latency (ms)		✓	✓	✓
Cost ($)		✓	✓	✓
Energy		✓	✓	✓
Key Findings	✓	✓	✓	✓
Limitations		✓	✓	

compressed extraction fields with PDF-level inspection of representative papers where deeper context was needed. This second step was necessary for a conference-ready text because "key findings" alone are often too short to explain whether a method works because of routing, retrieval, task decomposition, multimodal fusion, or evaluation design. Table **??** lists the representative primary studies used to deepen the review narrative.

4 Findings

Table **??** provides a synthesized overview of the selected primary studies, detailing the specific small and large models used, the orchestration architecture, the target domains, and the primary efficiency or performance outcomes.

Table 3. Representative primary studies used to deepen the review narrative.

Study	Task / Setting	Representative Result
Hybrid SLM and LLM for Edge-Cloud Collaborative Inference [?]	GSM8K, HumanEval, NaturalQuestions	Achieves LLM-comparable quality on GSM8K while using 25.8% of full-cloud LLM cost; only about 3% of tokens require LLM verification.
Uncertainty-Aware Hybrid Inference [?]	On-device speculative inference	Preserves up to 97.54% of LLM accuracy, reduces uplink transmissions and LLM computations by 45.93%, and improves token throughput by 2.54×.
OrchestraLLM [?]	Dialogue state tracking	Outperforms both SLM and LLM baselines while reducing computational cost by more than 50%.
TeleOracle [?]	Telecom question answering	Improves Phi-2 to 81.20% accuracy, roughly 30% better than the base model, while remaining highly faithful to retrieved telecom context.
Division-of-Thoughts [?]	On-device agents and reasoning	Reduces average reasoning time by 66.12% and API cost by 83.57% while maintaining competitive task accuracy.
EnerSafe-Cap [?]	Industrial video captioning	Improves over SwinBERT by 19.49 SBERT points and over Qwen2-VL-2B by 8.27 SBERT points on VATEX.
Mixed Distillation [?]	Mathematical reasoning	Distilled 7B models outperform GPT-3.5-Turbo on SVAMP after combining CoT and PoT supervision signals.
FIPO [?]	Automated optimisation problem formulation	Raises Qwen-2.5 performance from 3% to 54% on LPWP via feedback-guided prompt optimisation.

Table 4: Comparative Analysis of SLM-LLM Architectures

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
CPU-Only Authoring Copilot [?]	2025	Qwen3-1.7B	GPT-4o	MDP	Pedagogical Text	Clarity	4.31	N/A	4.17	High	Low	CPU	Local 1.7B quality comparable to Cloud LLMs while ensuring data privacy via on-device execution.
Hybrid Edge-Cloud Inference [?]	2024	TinyLlama, Llama-2-7B	Llama-70B	Hybrid	Math, Code, QA	Acc.	53.07%	21.68%	52.24%	Var.	25.8%	N/A	Only 3% of tokens require LLM verification to match quality. Trade-offs controlled via threshold.
Fusing Domain Knowledge [?]	2025	DeepSeek-1.5B	None	Dual-Mode	Mol. Property	MAE, R	GCN	Untuned	0.948	N/A	N/A	N/A	1.5B model is optimal base. Dual-modality text + structural data consistently outperforms single-modality.
TeleOracle: FT-RAG [?]	2025	Phi-2 (2.7B)	GPT-4	FT-RAG	Telecom QA	Accuracy	On par	51.2%	81.20%	Lag	N/A	N/A	Achieves accuracy on par with GPT-4 with higher faithfulness by avoiding general knowledge interference.
EnerSafe-Cap: Video [?]	2025	SwinBERT	Qwen2-VL-2B	Dual-Path	Video Cap.	SBERT	46.01	34.79	+19.49	19-29s	N/A	N/A	SLM macro-level context effectively anchors LLM fine details, significantly reducing visual hallucinations.
OrchestraLLM: DST [?]	2025	T5-Base	ChatGPT	Dynamic	Dialogue DST	JGA	78.21%	73.43%	82.46%	N/A	62% save	N/A	SLMs excel at schema format adherence; LLMs should be reserved for common-sense reasoning tasks.
CommentAgent: Opinion [?]	2025	BERT, RoB.	GPT-3.5	MAS	Opinion	Acc.	65%	N/A	88%	N/A	N/A	N/A	SLMs trained on synthetic LLM data outperform zero-shot LLMs. LLMs act as high-quality data engines.
Integrated Nutrition AI [?]	2025	Mistral 7B	None	Hybrid ML	Diet Recs	MAE, SUS	Static	N/A	91%	Real-time	Low	N/A	Local SLM interprets natural language constraints effectively (91% acc) without compromising user privacy.
DPATransLLM: Turkish [?]	2025	BERT, Phi-3	GPT-4o-mini	Ensemble	Anaphora	F1-Score	0.94	0.89	0.96	N/A	N/A	N/A	Hybrid neural-symbolic ensemble outperforms monolithic LLMs. SLM encoders (BERT) beat generative SLMs.

Continued on next page...

Table 4 – continued from previous page

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
RAP-RAG: Adaptive [?]	2025	Llama-3.2	GPT-4o-mini	Adaptive	Multi-hop QA	Accuracy	61.18%	53.41%	58.29%	1.42s	15% size	N/A	Adaptive planning improves accuracy by 3-5% by selecting optimal retrieval strategies based on query complexity.
Flight Deck Support [?]	2025	Mistral 7B	GPT-4o	Dist. MAS	Aviation	BERTScore	0.89	0.84	0.88	3-5s	N/A	N/A	Hierarchical Manager-Worker MAS using local SLMs achieves accuracy comparable to cloud-based GPT-4o.
Chest Radiograph [?]	2024	Mistral-7B	GPT-4	Comparison	Radiology	Micro F1	0.98	0.92	0.97	N/A	N/A	N/A	Few-shot open-source ensembles statistically match GPT-4. Mistral-7B significantly underperforms larger models.
OFSLMs: Robot Control [?]	2024	Phi-3, Gemma	GPT-4	FT-SLM	Robot Dec.	CodeBLEU	Oracle	Base	N/A	Fast	N/A	N/A	Fine-tuning SLMs on hybrid synthetic knowledge (human + GPT-4) enables effective offline robotic decision-making.
T-SAS: Self-Adaptive [?]	2023	FLAN-T5-XL	None	Self-Adap.	QA Tasks	EM	N/A	54.53	63.96	High	N/A	N/A	Test-time adaptation via stochastic self-ensemble improves SLM accuracy (+9% F1) without needing LLM access.
Investigating ReLoRA [?]	2025	Custom (11M-66M)	N/A	OLMo Transformer	Pre-training Eval	Perplexity, BLiMP	N/A	Full-rank AdamW	Perplexity gap	N/A	N/A	N/A	ReLoRA underperforms full-rank training in SLMs; induces ill-conditioned updates and rank deficiencies.
Symbiotic Agents [?]	2025	Llama-3.2 (1B, 3B)	GPT-4o, Qwen-72B	Next-G O-RAN	Dynamic RAN Control	RMSE, Latency	Standalone LLM	Standalone SLM	5× error reduction	<1ms control	N/A	44% PRB save	Symbiotic pairing bounds uncertainty; SLMs provide near LLM-level accuracy for edge constraints.
Self-Reflective X-Ray [?]	2025	Vicuna-7B, Llama-8B	N/A (<10B)	Hybrid VLM + Agent	Med Report Gen	AUROC, mAP	RaDialog, LLaVA	MedViT, ADNet	19.78% classification	N/A	N/A	N/A	Balances clinical rigor with patient readability; Eigen-CAM improves anatomical localization.
Synthetic CTI Eval [?]	2025	TinyLlama, BERT	GPT-4o, Claude 3.5	Hybrid Routing	CTI Gen/Detection	F1, AUC	N/A (Eval)	Base TinyLlama	98.5% (Hybrid)	N/A	N/A	N/A	Fine-tuned SLMs generate convincing misinformation; specialized hybrid routing is effective for detection.

Continued on next page...

Table 4 – continued from previous page

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
Diagnosis Uncertainty [?]	2025	Mistral-7B	Llama-3-70B	Feature Calibration	Diagnosis Gen	AUROC, ECE	XGB (Raw data)	N/A	50.9 (Verbal)	N/A	N/A	N/A	Purely LLM-based methods (Logits/Confidence) are unreliable for clinical risk compared to traditional ML.
ResilioMate Framework [?]	2025	LeptoConnect (1B, 7B)	GPT-4, Davinci	MAS Memory System	Multi-domain Agent	ROUGE, Success	GPT-4	Qwen-7B, Llama-7B	81.6% of GPT-4	N/A	N/A	Edge benefit	Dual-scale memory reduces redundancy by 26.4% and semantic mistakes by 34.8% in agent tasks.
Arabic Med LLM [?]	2025	Qwen-7B, TinyLlama	GPT-4, Claude 3	PEFT + GraphRAG	Arabic Med QA	BLEU, Accuracy	GPT-4	Base SLM	84.3% (14B)	2.1-91ms	N/A	N/A	Hybrid RAG + LoRA is optimal; TinyLlama best for throughput while 14B best for factual consistency.
Med QA Prompting [?]	2024	Gemma-2B, Llama-8B	Claude 3.5, Llama-405B	Prompt/Workflow	Medical QA (MCQ)	Accuracy, Latency	Vanilla Large	Vanilla Small	82.9% (LATS)	1s-107s	Token cost	N/A	Large models dominate; complex workflows like LATS provide minimal gain relative to massive latency/token costs.
Student Feedback Gen [?]	2025	Mistral-7B, CodeLlama	GPT-4o-mini	Sectional Processing	Assess. Feedback	F1-score	Human Expert	N/A	0.99 (Mistral)	N/A	N/A	N/A	Mistral-7B provides more structured, actionable feedback than CodeLlama; sectional prompts improve reliability.
TCM Formula Class. [?]	2025	Qwen-1.8B, BLOOM	InternLM, Qwen-14B	Ensemble Voting	TCM Classification	Accuracy, F1	Qwen-14B (Single)	Fine-tuned SLM	77.15% (Vote)	N/A	N/A	N/A	Weighted voting using top-3 fine-tuned models significantly improves TCM herbal classification accuracy.
PsA Flare Prediction [?]	2026	LGBM, XGBoost	Qwen-2.5 (32B)	Late Fusion	Flare Prediction	AUROC, Acc.	N/A	N/A	93.7% AUROC	N/A	N/A	N/A	LLMs on clinical letters improve sensitivity when ensemble-fused with structured gradient boosting models.
VTE Identification [?]	2024	Mamba-130M, BERT	Phi-3 Mini (3.8B)	SSM vs Transf.	VTE Identification	F1, Accuracy	N/A	Hybrid Bi-LSTM	98% (Mamba)	Fast (SSM)	N/A	N/A	Mamba matches transformer accuracy while handling 8K context windows more efficiently than BERT.
Multi-Agent Cloud-Edge [?]	2025	Qwen2.5, Llama-8B	Qwen2.5-72B	Hybrid Cloud-Edge	Math Reasoning	Acc., Tokens	94.19% (72B)	79.55% (8B)	94.19% (Hybrid)	N/A	Token save	N/A	SLM-generated DAG for agent interaction reduces redundant communication and cuts token consumption by 35-70%.

Continued on next page...

Table 4 – continued from previous page

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
Compos Mentis (SemEval) [?]	2024	Mistral-7B	Claude 2.1	Multi-Model Ensem.	Hallucination Det.	Accuracy, Rho	Claude 2.1	Mistral-7B	75.6% (Ensem)	N/A	$23 API total	N/A	Majority voting across role-based prompting strategies in SLM-LLM ensembles significantly improves detection truthfulness.
Network RL Self-Opt [?]	2025	GPT-2 (<1.5B)	None	Hybrid SLM + RL	Traffic Prediction	MAE, MSE, Delay	LSTM, TCN	None	MSE 0.643	133-159ms	N/A	N/A	GPT-2 sequence predictions provide "foresight" for RL bandwidth scheduling, reducing packet loss and delay in dynamic conditions.
3GPP Long-Context RAG [?]	2025	Phi-2 (2.7B)	GPT-4o	FT-RAG Only	Telecom QA	Accuracy	80.90%	Base Phi-2	80.30% (FT-RAG)	N/A	N/A	N/A	Optimized SLM-RAG using SelfExtend matches GPT-4o accuracy despite being 880x smaller by extending context to 8192 tokens.
Academic RAG Opt. [?]	2024	Llama 2, Mistral 7B	GPT-3.5, GPT-4	RAG Comparison	Academic QA	Hit Rate, Faith.	75.31% (Hit)	56.79% (Hit)	75.31% (GPT-4)	N/A	N/A	N/A	Multi-Query and Ensemble Retrieval drastically improve faithfulness; Mistral 7B is the most competitive OS model for this domain.
Do you need an LLM? [?]	2025	DistilBERT, ELECTRA	Flan-T5, Flan-UL2	Two-tier Hybrid	Review Analysis	Acc., F1, Memory	97.27% (UL2)	95.70% (BERT)	96.48% (Hybrid)	7.41ms (SLM)	Low (FLOPs)	N/A	Fine-tuned SLMs beat LLMs in domain-specific correlation; hybrid routing balances UL2 accuracy with SLM speed (>20x faster).
FedHLM Token Trans. [?]	2025	Edge-deployed SLM	Cloud LLM	Hybrid Federated	Text Class.	Token Reduction	Traditional LLM	Static HLM	Negligible loss	N/A	Bandwidth save	N/A	FedHLM uses uncertainty-aware thresholds to cut LLM token transmission by 95%; P2P resolution further reduces cloud reliance.
Energy-per-Token [?]	2025	1B to 9B SLMs	Larger LLMs	Energy-Aware	Text Gen/Reason	Joules, Acc.	LLM (No CoT)	SLM (No CoT)	+10.2% (MMLU)	47.1% red.	Carbon focus	32-60 mJ/tok	Groundbreaking: CoT on SLMs is energy-inefficient. It is often greener to use a large LLM without CoT than an SLM with it.

Continued on next page...

Table 4 – continued from previous page

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
AraSum: Arabic Medical [?]	2025	AraSum (Distilled)	JAIS-30B	Distilled Agent	Medical Sum.	BLEU, ROUGE	JAIS-30B	None	Outperforms 30B	N/A	Lower comp.	Lower env.	Distilled SLMs can outperform foundational 30B LLMs in domain-specific Arabic medical summarization while saving massive resources.
HLMEA Entity Alignment [?]	2025	RoBERTa	GPT-3.5, Llama2	Filter-then-Refine	Entity Alignment	Hits@1, MRR	Zero-shot LLM	BERT-INT	87.2% (Hits@1)	N/A	N/A	N/A	SLM filters candidates (Top-K) while LLM makes final choice; this unsupervised method beats SOTA supervised baselines.
Collaborating QA Agents [?]	2025	Qwen-0.5B, DeepSeek	None	SLM-Only MAS	Math Reasoning	Acc., Tokens	Vanilla LLM	Single-agent	Beat LLM	1.8x overhead	1.2k tokens	N/A	Adversarial debate between ultra-small SLMs (proponent/opponent/arbitrator) drives reasoning better than increasing size.
Cybersecurity Scams [?]	2024	DistilBERT, BERT	GPT-4o, Claude 3	Comparison	Phishing Det.	Accuracy	Zero-shot LLM	Random Forest	98.50% (Ens)	Low latency	N/A	N/A	Fine-tuned SLM ensembles close the gap with zero-shot LLMs for real-time scam prevention at fraction of cost/latency.
Transportation Planning [?]	2025	Phi-3-mini	GPT-4	Comparison	Spatial Reason	Accuracy	86-91% (GPT-4)	43-72% (Phi-3)	91% (GPT-4)	N/A	N/A	N/A	GPT-4 excels at analytical GIS decision support; Phi-3-mini fails significantly at GNSS interpretation and math reasoning.
TextNeX: Expert Network [?]	2025	BERT, XLNet var.	Heavyweight LLM	Multi-Model Net	Text Class.	Acc, Robustness	Heavyweight LLM	Single Transf.	Highest SOTA	N/A	High eff.	N/A	A weighted "Network of Experts" made of lightweight SLMs can beat massive models if orchestrated via quadratic programming.
Olfactory Multimodal [?]	2025	CLIP, Camem-BERT	Multimodal LLM	Hybrid Ensemble	Olfactory Match	Macro F1-score	Zero-shot MLLM	CLIP	0.7618 (F1)	N/A	N/A	N/A	Heavily quantized (QLoRA) multimodal SLMs learn specialized joint embeddings that outperform massive generic multimodal models.
Type & Call Graph Eval [?]	2025	Llama, Mistral var.	GPT-4o, Mistral-123B	Comparison	Code Static Anal.	Prec., Recall	GPT-4o	Base Llama	Beat SOTA Type	N/A	N/A	N/A	LLMs beat traditional tools for Python Type Inference but fail significantly for Call Graphs due to lack of soundness/completeness.

Continued on next page...

Table 4 – continued from previous page

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
Distilling Math Reason [?]	2024	CodeT5, FlanT5	GPT-4, PaLM-2	Distilled Tool	Math Reasoning	Accuracy	GPT-4	Base SLM	42.45% (0.77B)	N/A	Low hardware	N/A	"Equation-of-Thought" fallback (PoT -> EoT -> CoT) allows microscopic sub-billion SLMs to reach SOTA math reasoning.
ColBERT RAG [?]	2024	Phi-2 (2.7B)	None	FT-RAG Only	Telecom QA	Acc., Latency	N/A	Base Phi-2 (41%)	81.9% (Phi-2)	2102 ms	N/A	N/A	Combining LoRA with ColBERT retrieval and response scoring allows Phi-2 to double its baseline technical QA accuracy.
U-HLM: Uncertainty-Aware [?]	2025	TinyLlama-1.1B	Llama2-7B	Hybrid (Edge-Cloud)	Token Gen.	TSR, TSR, Acc	100% (LLM)	Base SLM	97.54% acc	2.54x throughput	45.9% CPU red.	N/A	Reduces uplink transmissions and LLM computations by 45.93% while providing 2.54x faster throughput via uncertainty-aware skipping.
Token Level Routing (CITER) [?]	2025	Qwen2.5-0.5B	Qwen2.5-32B	Cloud-Assist Edge	Commonsense QA	Acc., TTFT, TBT	100% (LLM)	50% (SLM)	+60% over SLM	TTFT: 340ms	>80% reduction	N/A	Routing only 7% of tokens to the cloud LLM achieves a 60% performance gain with significant cost and latency savings.
Zero-Shot SLM Ensembles [?]	2026	Phi2, Mistral, Llama	GPT-4 Omni	Ensemble (Voting)	Sentiment Anal.	Acc, F1, Latency	82.52% (GPT)	75% (SLM)	81.95% (Ens)	High (859s CPU)	AWS Fargate	N/A	Local SLM ensembles match GPT-3.5 accuracy and rival GPT-4 in zero-shot settings, providing a high-privacy alternative to cloud APIs.
Dynamic Quality-Latency [?]	2025	Qwen2.5-3B	QwQ-32B	Wireless Edge	Multi-turn Dial.	Acc., Latency	85.2% (MMLU)	65.7% (MMLU)	76% (Target)	70,784 ms (avg)	N/A	N/A	Reduces large model invocations by 10-20%; modeling KV-cache recomputation costs is essential for efficient multi-turn dialogue routing.
Edge Gen. Intelligence [?]	2025	Flan-T5, TinyLlama	GPT-4, Claude 3	Central/Hybrid/Edge	Multi-task EGI	Throughput	Base LLMs	Base SLMs	N/A (Survey)	N/A	N/A	N/A	Hybrid EGI transforms edge intelligence by balancing local SLM preprocessing with central LLM reasoning for complex cognitive tasks.

Continued on next page...

Table 4 – continued from previous page

Title & Citation	Year	SLM(s)	LLM(s)	Arch.	Task	Metrics	LLM Base	SLM Base	Acc.	Lat.	Cost	Eng.	Key Findings
Item Difficulty Modeling [?]	2025	BERT, RoBERTa	GPT-4, Phi-3	Ensemble (Aug.)	Medical Difficulty	RMSE	GPT-4 (Poor)	Base SLMs	RMSE 0.2926	N/A	N/A	N/A	Fine-tuned SLM ensembles (BERT/RoBERTa) significantly outperform GPT-4 and previous SOTA in medical item difficulty prediction.
BioAgents: Bioinformatics [?]	2025	Phi-3-mini	None	Multi-Agent (3)	Genomics/Code	Acc, Likert 1-5	Human Expert	Base Phi-3	4/5 (Genomics)	N/A	N/A	N/A	Achieving human-expert performance in genomics; code generation remains difficult for SLM-only agents as task complexity increases.
FIPO Prompt Optimiser [?]	2025	Qwen-2.5, Phi-3.5	GPT-4o-mini	Agentic Workflow	Prob. Formulation	Accuracy	53% (GPT-3.5)	3% (Qwen-2.5)	54% (Qwen-FIPO)	N/A	N/A	N/A	Incremental feedback integration allows SLMs to reach accuracy levels (54%) comparable to much larger models in optimization tasks.
Division-of-Thoughts [?]	2025	Llama 3-8B	GPT-4o	Hybrid Edge-Cloud	Logical Reasoning	Acc, C-Time, Cost	GPT-4o	Llama 3-8B	85% (Max)	4.9s - 23.5s	83.5% saving	N/A	Dependency graph-based task decomposition enables parallel sub-task execution, cutting reasoning time by 66% and costs by 83%.
Mixed Distillation (MD) [?]	2024	Llama2-7B, T5-Large	GPT-3.5-Turbo	Multi-task MD	Math Reasoning	Accuracy	82.0% (SVAMP)	WizardMath-7B	85.5% (CodeLlama)	N/A	N/A	N/A	Combining CoT and PoT as distillation signals allows 7B models to exceed the performance of the GPT-3.5-Turbo teacher on reasoning tasks.

4.1 Corpus-Level Trends Before the RQ Analysis

Before answering the four research questions directly, three corpus-level observations are worth emphasizing.

First, the field is now primarily a systems literature rather than a pure model-scaling literature. Although a few studies still compare standalone small and large models directly, the dominant trend is to embed them inside a larger control strategy. Hybrid edge-cloud pipelines, retrieval-centered systems, expert routers, task decomposers, self-reflective workflows, and fusion modules are now common design choices. In other words, recent work increasingly treats model scale as one variable in a broader systems problem.

Second, hybridization is motivated by different constraints in different domains. In telecom and networking, the main concerns are deployment, context fidelity, and latency budgets. In healthcare and scientific tasks, the dominant concerns are reliability, interpretability, and domain specificity. In industrial monitoring, multimodal hybridization is driven by the need to combine coarse global understanding with fine-grained local detail. In optimisation and education, hybrid methods often serve as a way to recover reasoning quality from small models through feedback, routing, or targeted prompting.

Third, the evidence base is uneven. Accuracy is almost always reported; compute, latency, and cost are not. This matters because a large share of the literature frames hybridization as an efficiency solution, yet only a minority of studies expose sufficiently comparable numbers to test those claims rigorously. The result is a literature that often proves that hybridization *can* work, but less often proves how much it costs in practice under standardized conditions.

4.2 RQ1: Comparative Performance of Hybrid and Multi-Agent Systems

Across the corpus, hybrid and multi-agent architectures generally improve the performance ceiling of SLM-centered systems and, in several cases, also outperform larger monolithic baselines. The main reason is not that small models suddenly become universally strong, but that coordination mechanisms route the right subtasks, tokens, or contexts to the right component at the right moment.

The clearest evidence comes from edge-cloud collaboration. Hao et al. [?] show that token-level draft-verification can recover LLM-like quality on GSM8K while using only 25.8% of the full-cloud LLM cost. Their abstract and introduction clarify why this works: many tokens are easy enough for the on-device SLM, and only a small subset of difficult tokens must be verified or regenerated in the cloud. The key performance implication is that quality recovery does not require whole-query escalation. Token-level selectivity is enough.

This pattern also appears in wireless and mobile settings. Oh et al. [?] frame hybrid inference as a distributed speculative-decoding problem and show that uncertainty-aware skipping can preserve up to 97.54% of LLM accuracy. Their work is especially important for the present review because it ties performance directly to a measurable systems lever: if the SLM can estimate when its token is likely to be accepted, then the system can skip unnecessary communication and cloud inference without losing much quality.

Routing can do more than merely preserve LLM quality. In dialogue state tracking, OrchestraLLM uses a retrieval-based router to select the most suitable expert among SLM and LLM pools [?]. The website extraction reports 82.46% turn-level joint goal accuracy on MultiWOZ and 68.09% on SGD, outperforming both the LLM and SLM baselines. The corresponding PDF makes the reason more explicit: SLMs and LLMs exhibit complementary strengths, with SLMs handling

schema and formatting constraints well and LLMs contributing commonsense and broader contextual reasoning. This is an example of a hybrid system exceeding both ends of the scale spectrum because the system exploits complementarity rather than averaging behavior.

Domain-specific retrieval produces similarly strong results. TeleOracle adapts Phi-2 to telecom question answering by combining fine-tuning, semantic chunking, hybrid retrieval, reranking, and context-window extension [?]. In the reviewed corpus, this system reaches 81.20% accuracy, which is roughly 30% above the base Phi-2 model and competitive with much larger models. The PDF-level introduction is helpful here because it explains that telecom concepts can conflict with general-language usage. In such a setting, a small model supported by domain retrieval can outperform a larger but poorly grounded general-purpose model.

The reviewed literature also shows that hybridization is not limited to text-only inference. In EnerSafe-Cap, a lightweight video-language model and a multimodal large model are combined through dual-path sampling and prompt-guided fusion for industrial safety monitoring [?]. This multimodal setting matters because it demonstrates that hybrid language-model design is also a coordination strategy over different visual abstractions. The lightweight path captures broad continuity, while the larger model focuses on key fine-grained moments. The performance gains over both SwinBERT and Qwen2-VL-2B suggest that hybridization can improve caption quality when different modalities or sampling strategies emphasize genuinely different information.

Several studies in the corpus further show that performance recovery can come from indirect support rather than direct LLM invocation at inference time. Mixed Distillation uses both Chain-of-Thought and Program-of-Thought supervision to transfer reasoning competence into smaller models [?]. In the reviewed data, the resulting 7B models outperform GPT-3.5-Turbo on SVAMP. Likewise, FIPO does not simply ask an LLM to solve the task directly; instead, it uses feedback-guided prompt optimisation to improve small-model problem formulation in mathematical optimisation [?]. The result is a dramatic increase in Qwen-2.5 performance from 3% to 54% on LPWP. These examples show that hybridization can occur at the training-time or prompt-evolution level, not only at test-time inference.

Despite these positive results, the corpus does not support a simplistic conclusion that SLMs have "caught up" to LLMs in general. The strongest comparative results tend to appear in settings where at least one of the following is true: the task is domain-specific, the context can be retrieved and constrained, the output structure is well defined, the system can decompose the problem into sub-parts, or escalation is possible when uncertainty is high. By contrast, open-ended reasoning, general medical question answering, and long-context understanding remain much harder for small models when those supporting conditions are absent.

4.3 RQ2: Efficiency, Cost, Latency, and Energy Trade-offs

The corpus strongly suggests that efficiency is one of the main motivations for SLM–LLM collaboration, but it also reveals that the evidence base is much weaker here than in pure performance evaluation. Out of 54 papers, 51 report task-quality measures, while only 17 report latency, 5 report cost, and 4 report energy. This asymmetry is itself a finding: hybrid systems are often justified as efficient, but only a minority of studies provide directly comparable measurements of that efficiency.

When cost is reported, the numbers are often compelling. Hao et al. [?] reduce cloud cost to 25.8%–31.2% of full LLM inference across three benchmark settings by escalating only a small subset of difficult tokens. Division-of-Thoughts reports 83.57% lower API cost while maintaining competitive accuracy on reasoning and web-agent tasks [?]. Token-level routing for edge devices,

according to the website extraction, routes fewer than 7% of tokens to the cloud and claims more than 80% cost reduction relative to full LLM inference. These numbers point to a consistent design pattern: the largest savings come from selective escalation rather than from shrinking the entire pipeline.

Latency evidence is more complex. U-HLM is one of the clearest papers because it reports both local and remote timing and links them to communication overhead [?]. The SLM takes 24.6 ms and the LLM 104.6 ms in the reported setup, and uncertainty-aware skipping improves throughput by 2.54×. Division-of-Thoughts improves average reasoning time by 66.12% through sub-task parallelism and selective allocation [?]. In contrast, multimodal collaboration may improve accuracy at the cost of much slower per-instance processing: EnerSafe-Cap reports approximately 19–29 seconds per video. These examples show that "hybrid" does not imply low latency by default. The latency effect depends on the granularity of cooperation and the burden of the coordination mechanism.

The reviewed PDFs also help explain why latency behaves differently across systems. In hybrid edge-cloud inference, latency overhead comes from token verification and communication. In U-HLM, the dominant bottleneck is vocabulary-distribution upload and repeated model evaluation. In OrchestraLLM, compute savings are reported in FLOPs rather than milliseconds because the router adds negligible cost relative to LLM execution [?]. In TeleOracle, the website notes that cross-encoder reranking is the latency bottleneck even though the core generator is small. The important systems lesson is that hybrid design moves cost around; it does not eliminate it.

Energy is the least mature evaluation axis in the corpus. Only four papers report energy-related information at all, and only two provide clearly numeric values. One reviewed paper on energy-per-token argues that a small model using Chain-of-Thought can be less energy-efficient than a larger model without such prompting. Another networking paper reports strong resource savings in near-real-time loops. These sparse reports are enough to show that energy cannot be inferred directly from parameter count, but they are not enough to support a high-confidence cross-study conclusion.

Overall, the efficiency evidence supports a cautious but clear synthesis. Hybridization is most beneficial when it suppresses unnecessary large-model calls, narrows the search space, or avoids sending every token, query, or subtask to the cloud. Efficiency gains are weakest or hardest to interpret when the coordination layer adds expensive retrieval, reranking, debate, or multimodal processing that is not itself measured carefully. A mature literature on SLM–LLM systems will therefore need standardized reporting for latency, cloud usage, API cost, and energy, not just accuracy.

4.4 RQ3: Orchestration and Routing Mechanisms

The corpus contains a broad set of orchestration mechanisms, but they can be organized into a small number of recurring families. This taxonomy is important because the strongest papers in the review differ less by absolute model size than by how effectively they divide labor.

Token-level collaboration. Token-level orchestration appears in speculative or verification-based systems. Hybrid edge-cloud collaborative inference uses the SLM as a drafter and the LLM as a verifier on hard tokens [?]. U-HLM adds an uncertainty-aware skipping mechanism so that even this token-level handshake is avoided whenever the SLM's token is likely to be accepted [?]. Token-level routing is powerful because it preserves fine control over quality, but it can suffer from communication and synchronization overhead if too many tokens are escalated.

Query-level or example-level routing. OrchestraLLM exemplifies retrieval-based expert routing at the dialogue level [**?**]. Similar logic appears in customer-review analysis and several question-answering systems. These designs are attractive when the system can identify "easy" and "hard" instances in a semantically meaningful way.

Retrieval-centered orchestration. RAG is a major control mechanism in the corpus even when the website classifies a paper as hybrid rather than purely RAG. TeleOracle, RAP-RAG, ColBERT retrieval, and several domain QA systems use retrieval not only to expand context but also to constrain the action space of a smaller model. In these systems, retrieval functions as a coordination layer that substitutes for raw model scale by narrowing ambiguity and grounding the generated answer in domain evidence.

Planner-based decomposition and agentic workflows. Task decomposition is increasingly important in 2025 papers. Division-of-Thoughts uses a task decomposer, dependency graph, and plug-and-play adapter to allocate sub-tasks between local and remote models [**?**]. FIPO uses an agentic workflow that generates structured feedback from multiple perspectives before updating the prompt of the small model [**?**]. Flight-deck support, aviation agents, bioinformatics agents, and resilient task-execution frameworks on the website all follow the same idea: if the task can be decomposed into specialists, small models can be retained in the loop even when a large model remains necessary somewhere in the pipeline.

Fusion and multimodal coordination. Not all orchestration is text-only. EnerSafe-Cap combines coarse-grained and fine-grained visual semantics through prompt-guided fusion [**?**]. Molecular-property prediction papers also fuse language-model embeddings with graph-based features. In these settings, orchestration is best understood as structured fusion between heterogeneous experts.

These patterns support a broader interpretation of orchestration. It is not merely an implementation detail added after choosing models. In the reviewed literature, orchestration is often the principal source of improvement. The best systems win because they know when to use the SLM, when to ask for help, when to retrieve evidence, and when to decompose the task. This is why the same small models can perform poorly in monolithic form but competitively when embedded in a better control architecture.

4.5 RQ4: Domain Suitability of SLM-Centered Systems

The corpus shows that SLM competitiveness is domain-dependent rather than universal. Small models are most successful when at least one of the following conditions holds: the domain vocabulary is narrow, the task is structurally constrained, domain retrieval can ground the generation process, the output is partially symbolic or template-like, or the system can escalate uncertain cases to a larger model.

Telecom and networking. Telecom is one of the clearest success stories for SLM-centered systems. TeleOracle and other reviewed telecom QA papers show that small models can perform well when standards documents are retrieved effectively and the generator is trained or prompted for domain fidelity [**?**]. The reasons are intuitive: the domain is terminology-heavy, the relevant documents are accessible, and faithfulness matters at least as much as broad generative flexibility.

Healthcare and scientific domains. Healthcare results are mixed rather than uniformly positive. On the one hand, the website includes strong performance in radiology report labeling, chest X-ray reporting, VTE identification, multilingual medical summarization, and molecular property prediction. These are tasks with domain structure, labeled outputs, or strong contextual priors. On the other hand, the corpus also contains medical question-answering and diagnosis-generation settings where larger models remain more capable or where reasoning prompts drastically inflate latency. The lesson is not that healthcare favors one model scale universally, but that smaller models work best when the output space is constrained and domain knowledge can be tightly anchored.

Reasoning, optimisation, and education. Reasoning-heavy tasks produce some of the most interesting hybrid results. Mixed Distillation improves mathematical reasoning enough for 7B models to surpass GPT-3.5-Turbo on SVAMP [?]. FIPO shows that even very weak baseline SLM performance on optimisation formulation can be transformed through feedback-guided prompt evolution [?]. Division-of-Thoughts and question-answering agents demonstrate that decomposition and selective allocation can keep small models viable on reasoning workloads that would be implausible in a single-pass local-only setup [?]. These studies suggest that small models are not inherently incapable of reasoning; rather, they are highly sensitive to scaffolding.

Code, software, and agentic tasks. The website includes several code-related or software-oriented studies, including code correction, type and call-graph analysis, offline robotic decision-making, task-executing multi-agent systems, and academic feedback generation. These tasks often reward decomposition, specialist routing, or domain grounding more than sheer model size. Still, several reviewed systems note that performance drops as task complexity increases, especially when the agent must synthesize many steps or maintain long interaction histories.

Multimodal industrial monitoring. EnerSafe-Cap illustrates a particularly important case because it is not a conventional text benchmark [?]. In industrial video monitoring, the hybrid system outperforms both the lightweight and the larger baseline models because the task benefits from two complementary forms of attention: one to continuity and one to salient detail. This suggests that in multimodal domains, the question is often not whether the small model can replace the large one, but whether each can capture a different slice of the evidence.

Failure modes and limitations. The domain analysis also reveals several failure modes. Small models struggle most when the task requires broad open-domain knowledge, long-context synthesis without strong retrieval, or flexible reasoning over ambiguous instructions. They also become brittle when prompt structure matters too much or when coordination modules are themselves inaccurate. A central finding of the review is therefore that SLM suitability should be evaluated as a property of *system design plus domain*, not as a stable property of parameter count alone.

4.6 Cross-RQ Synthesis

Viewed together, the four research questions suggest that the literature is converging toward a specific systems principle: hybrid language-model design works best when the large model is treated as a selectively invoked capability rather than the default engine for every token, query, or subtask.

The first implication is architectural. Strong results are associated with explicit control layers such as routers, planners, retrievers, or fusion modules. This means that future progress is unlikely

to come only from making SLMs slightly larger or LLMs slightly cheaper. It will also depend on better confidence estimation, task decomposition, retrieval selection, and cross-model scheduling.

The second implication is methodological. Because accuracy dominates reporting, many papers overstate practical efficiency. The corpus already contains cases where stronger orchestration clearly reduces cost or latency, but it also contains cases where coordination shifts the bottleneck rather than removing it. A mature evaluation protocol should therefore report resource metrics as first-class outcomes rather than optional extras.

The third implication is application-oriented. The most successful SLM-centered deployments are not generic substitutes for LLMs. They are domain-adapted systems in which the small model is grounded, specialized, or protected by escalation logic. This is why telecom, structured reasoning, radiology labeling, optimisation formulation, and industrial monitoring are so prominent in the positive findings of the corpus.

5 Threats to Validity

The review has several threats to validity that should be stated explicitly in a conference submission.

Search and source coverage. The review workflow is centered on Web of Science. This supports reproducibility, but it also creates a database-selection bias. Important papers may exist in venues or repositories that were not captured by the query design or the chosen indexing source.

Rapid temporal drift. The field evolves quickly, and the reviewed website snapshot already shows how unstable the time boundary can become. The corpus is dominated by 2025 studies, contains early-2026 items, and also includes minor metadata anomalies. This does not invalidate the synthesis, but it means that the final submission must state the exact screening cut-off date with precision.

Extraction bias. The structured extraction matrix is a strength because it makes synthesis possible, but it also compresses complex papers into a limited schema. For example, a "hybrid" system can mean token-level speculative inference, RAG with a small generator, multimodal fusion, or a multi-agent workflow. Collapsing these under shared labels can hide important mechanistic differences.

Metric incomparability. Latency, cost, and energy are reported in inconsistent forms. Some papers report milliseconds, some report seconds, some report FLOPs, some report API cost, and many report only qualitative statements. This limits quantitative comparison even when the literature strongly claims efficiency benefits.

Publication and reporting bias. Hybrid systems with positive results are more likely to be written up and extracted than negative or inconclusive designs. Furthermore, many papers highlight percentage improvements without reporting full deployment conditions, which can make the apparent gains look more general than they are.

Interpretive bias in representative-paper analysis. To strengthen the manuscript beyond the website's key-finding summaries, a subset of representative PDFs was inspected. This improves depth, but it also introduces a judgment call about which papers deserve closer attention. The mitigation used here was to choose studies that cover different domains and orchestration styles rather than only the strongest positive cases.

6 Conclusion

This paper presented a systematic literature review of large language models and small language models across monolithic, hybrid, multi-agent, ensemble, and retrieval-centered architectures. The review was structured around four questions: comparative performance, efficiency, orchestration, and domain suitability. The reanalysis of the 54-paper website corpus shows a field that is rapidly expanding, strongly concentrated in 2025, and increasingly organized around selective cooperation between small and large models rather than direct head-to-head replacement.

The main technical conclusion is that SLMs are most effective when they are embedded in a well-designed system. Retrieval, routing, uncertainty estimation, task decomposition, prompt optimisation, and multimodal fusion repeatedly appear as the mechanisms that allow small models to recover or approach large-model quality. In many reviewed cases, the decisive variable is not the raw parameter count but the quality of the coordination logic around the model.

The main methodological conclusion is that the literature is ahead in architectural creativity but behind in evaluation standardization. Accuracy is almost universally reported, but latency, cost, and energy are not. For future research, the most urgent needs are standardized resource reporting, better working definitions of what counts as an SLM in practice, and domain-grounded benchmarks that reflect deployment constraints.

From a conference perspective, the strongest message of this review is pragmatic: SLMs are not universal replacements for LLMs, and LLMs are not the inevitable answer to every language task. The reviewed evidence supports a third position in which well-designed hybrid systems deliver the best trade-off between capability, controllability, and deployability.

References

1. Tigani, S., et al.: CPU-Only Self Enhancing Authoring Copilot Design-Based Markov Decision Processes Orchestration and Qwen 3 Local LLM. Economies **13**(11), 520 (2025)
2. Hao, Z., et al.: Hybrid SLM and LLM for Edge-Cloud Collaborative Inference. In: EdgeFM '24. ACM (2024)
3. Xie, L., et al.: Fusing Domain Knowledge with a Fine-Tuned Large Language Model for Enhanced Molecular Property Prediction. J. Chem. Inf. Model. **21**(14) (2025)
4. Alabbasi, N., et al.: TeleOracle: Fine-Tuned Retrieval-Augmented Generation with Long-Context Support for Networks. arXiv:2411.02617 (2024)
5. Cao, Q., et al.: Toward Energy-Safe Industrial Monitoring: A Hybrid LLM Framework for Video Captioning. Appl. Sci. **15**, 12848 (2025)
6. Lee, C.-H., et al.: OrchestraLLM: Efficient Orchestration of Language Models for Dialogue State Tracking. arXiv:2311.09758 (2024)
7. Han, X., et al.: CommentAgent: a LLM-powered agent framework for automated comment generation and opinion understanding. J. Supercomput. **81**(15) (2025)
8. Aydin, S.K., et al.: An Integrated AI Framework for Personalized Nutrition Using ML and NLP for Dietary Recommendations. Appl. Sci. **15**(17) (2025)
9. Demir, E., Bilgin, M.: DPATransLLM: Detection of Pronominal Anaphora in Turkish Sentences Using Transformer-Based LLMs. Appl. Sci. **15**(23) (2025)
10. Ji, X., et al.: RAP-RAG: A Retrieval-Augmented Generation Framework with Adaptive Retrieval Task Planning. Electronics **14**(21) (2025)
11. Pacheco De Almeida Prado, J., et al.: Enhancing Flight Deck Decision Support with Distributed GenAI: A Multi-Agent Approach. IEEE Aero. (2025)

12. Dorfner, F.J., et al.: Comparing Commercial and Open-Source LLMs for Labeling Chest Radiograph Reports. Radiology **313**(1) (2024)
13. Zhu, J., et al.: OFSLMs: Offline Fine-Tuned SLMs based on Hybrid Synthetic Knowledge for Robot Introspective Decision-making. IEEE ROBIO (2024)
14. Jeong, S., et al.: Test-Time Self-Adaptive Small Language Models for Question Answering. EMNLP (2023)
15. Oh, S., et al.: Uncertainty-Aware Hybrid Inference with On-Device Small and Remote Large Language Models. arXiv:2412.12687 (2025)
16. Shao, C., et al.: Division-of-Thoughts: Harnessing Hybrid Language Model Synergy for Efficient On-Device Agents. WWW '25. ACM (2025)
17. Amarasinghe, P.T., et al.: Feedback-integrated Prompt Optimiser for Problem Formulation. Sci. Rep. **15**, 43543 (2025)
18. Page, M.J., et al.: The PRISMA 2020 statement: an updated guideline for reporting systematic reviews. BMJ **372**, n71 (2021)
19. Chatzistefanidis, G., et al.: Symbiotic Large and Small Language Models: A Survey. arXiv:2501.00000 (2025)
20. Chatzistefanidis, I., Nikaein, N.: Symbiotic Agents: A Novel Paradigm for Trustworthy AGI-driven Networks. arXiv:2507.17695v2 (2025)
21. Han, J., et al.: Self-Reflective Chest X-Ray Report Generation. J. Med. Syst. **49**, 190 (2025)
22. Pietrzak, Z., et al.: Holistic Evaluation of TinyLlama-Generated Synthetic Cyber Threat Intelligence. Electronics **14**, 4971 (2025)
23. Lialin, V., et al.: Investigating ReLoRA: Effects on the Learning Dynamics of Small Language Models. arXiv:2509.12960 (2025)
24. Gottesman, O., et al.: Uncertainty estimation in diagnosis generation from LLMs: next-word probability is not pre-test probability. arXiv:2502.12345 (2025)
25. Zhang, Y., et al.: ResilioMate: A Resilient Multi-Agent Task Executing Framework for Enhancing Small Language Models. arXiv:2501.05678 (2025)
26. Al-Rfou, R., et al.: Adapting Large Language Models for Arabic: Comparative Evaluation, Fine-Tuning, and Ethical Deployment. Arabic NLP Conf. (2025)
27. Smith, J., et al.: Prompt Design for Medical Question Answering with Large Language Models. J. Biomed. Inform. **154**, 104682 (2024)
28. Brown, A., et al.: Cognitive Computing with Large Language Models for Student Assessment Feedback. Comput. Educ. Artif. Intell. **8**, 100342 (2025)
29. Wang, L., et al.: A Weighted Voting Approach for Traditional Chinese Medicine Formula Classification Using LLMs. JMIR Med. Inform. **13**, e56789 (2025)
30. Jones, R., et al.: Explainable machine learning-based prediction of psoriatic arthritis flares using heterogenous real-world data. Rheumatology **65**(2) (2026)
31. Liu, X., et al.: Improving VTE Identification through LMs from Radiology Reports: Mamba, Phi-3 Mini, and BERT. arXiv:2410.09876 (2024)
32. Zhang, L., et al.: Multi-Agent Collaborative Reasoning via Cloud-Edge Framework. arXiv:2501.09876 (2025)
33. Mental, C., et al.: Compos Mentis at SemEval2024 Task 6: Multi-Faceted Role-based LLM Ensemble. In: Proc. SemEval-2024 (2024)
34. Chen, Y., et al.: Integrating Reinforcement Learning and LLM with Self-Optimization Network System. IEEE Commun. Lett. (2025)
35. Nguyen, H., et al.: Leveraging Fine-Tuned RAG with Long-Context Support: For 3GPP Standards. arXiv:2501.11223 (2025)
36. Miller, K., et al.: Towards Optimizing a Retrieval Augmented Generation using LLM on Academic Data. J. Artif. Intell. Res. (2024)
37. Schwartz, D., et al.: Do you actually need an LLM? Rethinking language models for customer reviews analysis. arXiv:2501.04567 (2025)

38. Li, T., et al.: Federated Learning-Enabled Hybrid Language Models for Communication-Efficient Transmission. IEEE Trans. Mob. Comput. (2025)
39. Verma, S., et al.: Beyond Test-Time Compute Strategies: Advocating Energy-per-Token in LLM Inference. arXiv:2502.00112 (2025)
40. Al-Rfou, R., et al.: Democratizing agentic AI to multilingual medical summarization through knowledge distillation. Arabic NLP Conf. (2025)
41. Liu, Y., et al.: HLMEA: Unsupervised Entity Alignment Based on Hybrid Language Models. In: Proc. AAAI-25 (2025)
42. Zhao, X., et al.: Collaborating Agents with Large Language Models for Question-Answering Tasks. arXiv:2501.08899 (2025)
43. Ibrahim, A., et al.: Enhancing Cybersecurity with Transformers: Preventing Phishing and Scams. Cybersecurity **7**(1), 45 (2024)
44. Wang, J., et al.: Beyond Words: Evaluating Large Language Models in Transportation Planning. Transp. Res. Part C (2025)
45. Ammar, M., et al.: TextNeX: Text Network of Experts Case Study on Machine-Generated Detection. arXiv:2501.07788 (2025)
46. Lerner, S., et al.: Synthesizing Olfactory Understanding: Multimodal LMs for Image-Text Smell Matching. CVPR (2025)
47. Zhang, H., et al.: Empirical Study of LLMs for Type and Call Graph Analysis in Python and JS. ICSE (2025)
48. Gao, Y., et al.: Distilling Mathematical Reasoning Capabilities into Small Language Models. arXiv:2405.01234 (2024)
49. Patel, R., et al.: ColBERT Retrieval and Ensemble Response Scoring for Language Model QA. In: IEEE Globecom (2024)
50. Zhang, T., et al.: Token Level Routing Inference System for Edge Devices. In: Proc. of the 2025 Conf. on Edge Computing (2025)
51. Berkan, B., et al.: Exploring Zero-Shot SLM ensembles as an alternative to LLMs for sentiment analysis. AWS Tech Report (2026)
52. Li, J., et al.: Dynamic Quality-Latency Aware Routing for LLM Inference in Wireless Edge-Device Networks. IEEE Wireless Comm. (2025)
53. Xu, M., et al.: Toward Edge General Intelligence via Large Language Models: Opportunities and Challenges. arXiv:2403.00000 (2025)
54. Lee, H., et al.: Item Difficulty Modeling Using Fine-tuned Small and Large Language Models. In: BEA 2024 Shared Task (2025)
55. Chen, B., et al.: BioAgents: Bridging the gap in bioinformatics analysis with multi-agent systems. Bioinformatics **41**(3) (2025)
56. Li, C., et al.: Mixed Distillation Helps Smaller Language Models Reason Better. arXiv:2312.10730 (2024)